Ethics for Nerds

An Advanced Course in Computer Science
Summer Semester 2020

Precise Thinking 5.1
Discourse

Attacking an Argument

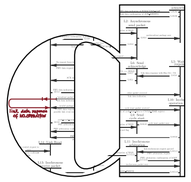

Prof. Holger Hermanns,
Kevin Baum, Sarah Sterz

UNIVERSITÄT
DES
SAARLANDES

PRECISE THINKING

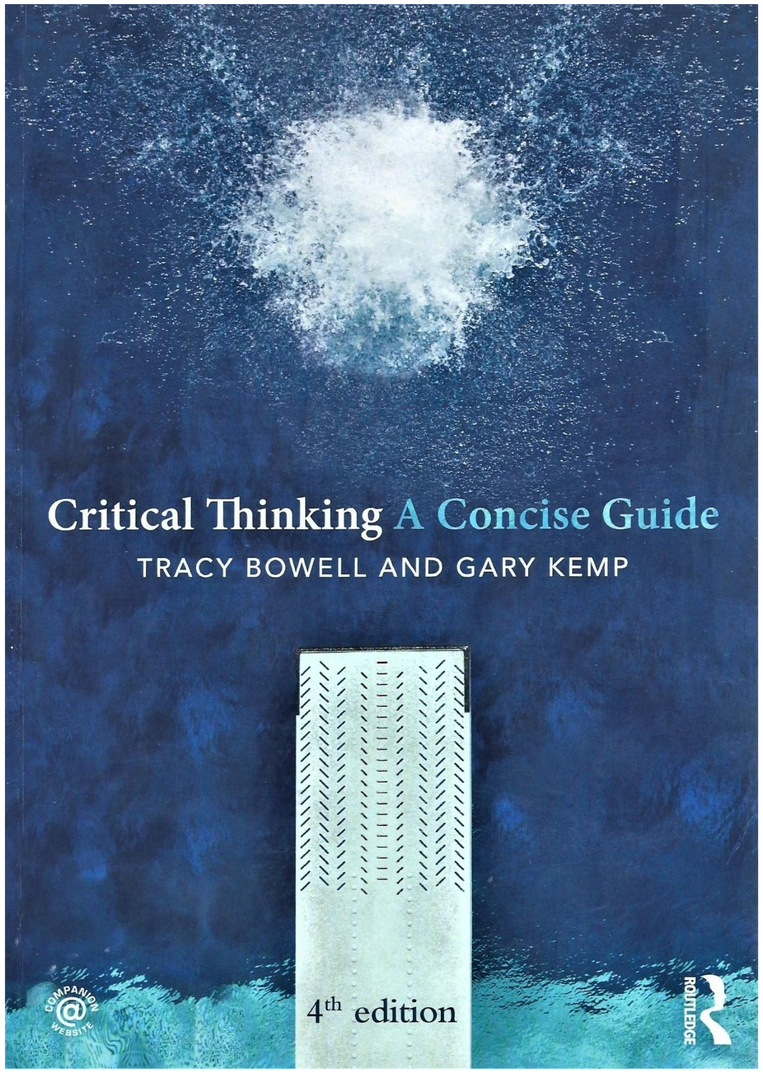

We *very* loosely follow this book:

> Bowell, T., & Kemp, G. (2015). *Critical thinking: A concise guide (4th ed.)*. Routledge.

However, *lots* of things deviate from the book. What is said in the lecture has precedence over the book.

You do **not** have to buy the book. If you want to have an inexpensive look, you can find a digital copy of the 2nd edition online at:

http://www.academia.edu/download/46383480/_Tracy_Bowell__Critical_Thinking_A_Concise_Guide_BookFi.org.pdf

Discourse is more than attacking arguments!

But attacking arguments is part of discourse.

ATTACKING AN ARGUMENT

Security through obscurity is the way to go – even for critical systems! If an attacker does not know where to attack, he cannot attack. So, any obscure system is a secure system and you do not really need more than that.

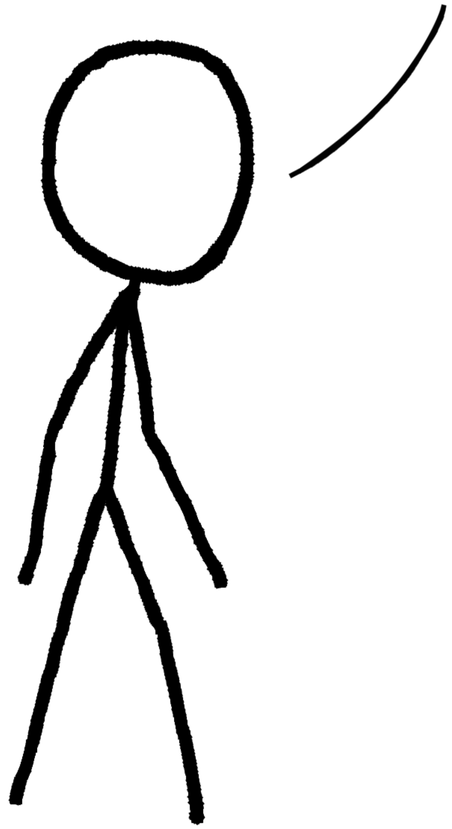

Sometimes, you disagree with arguments.
But how to attack them properly?

ATTACKING AN ARGUMENT

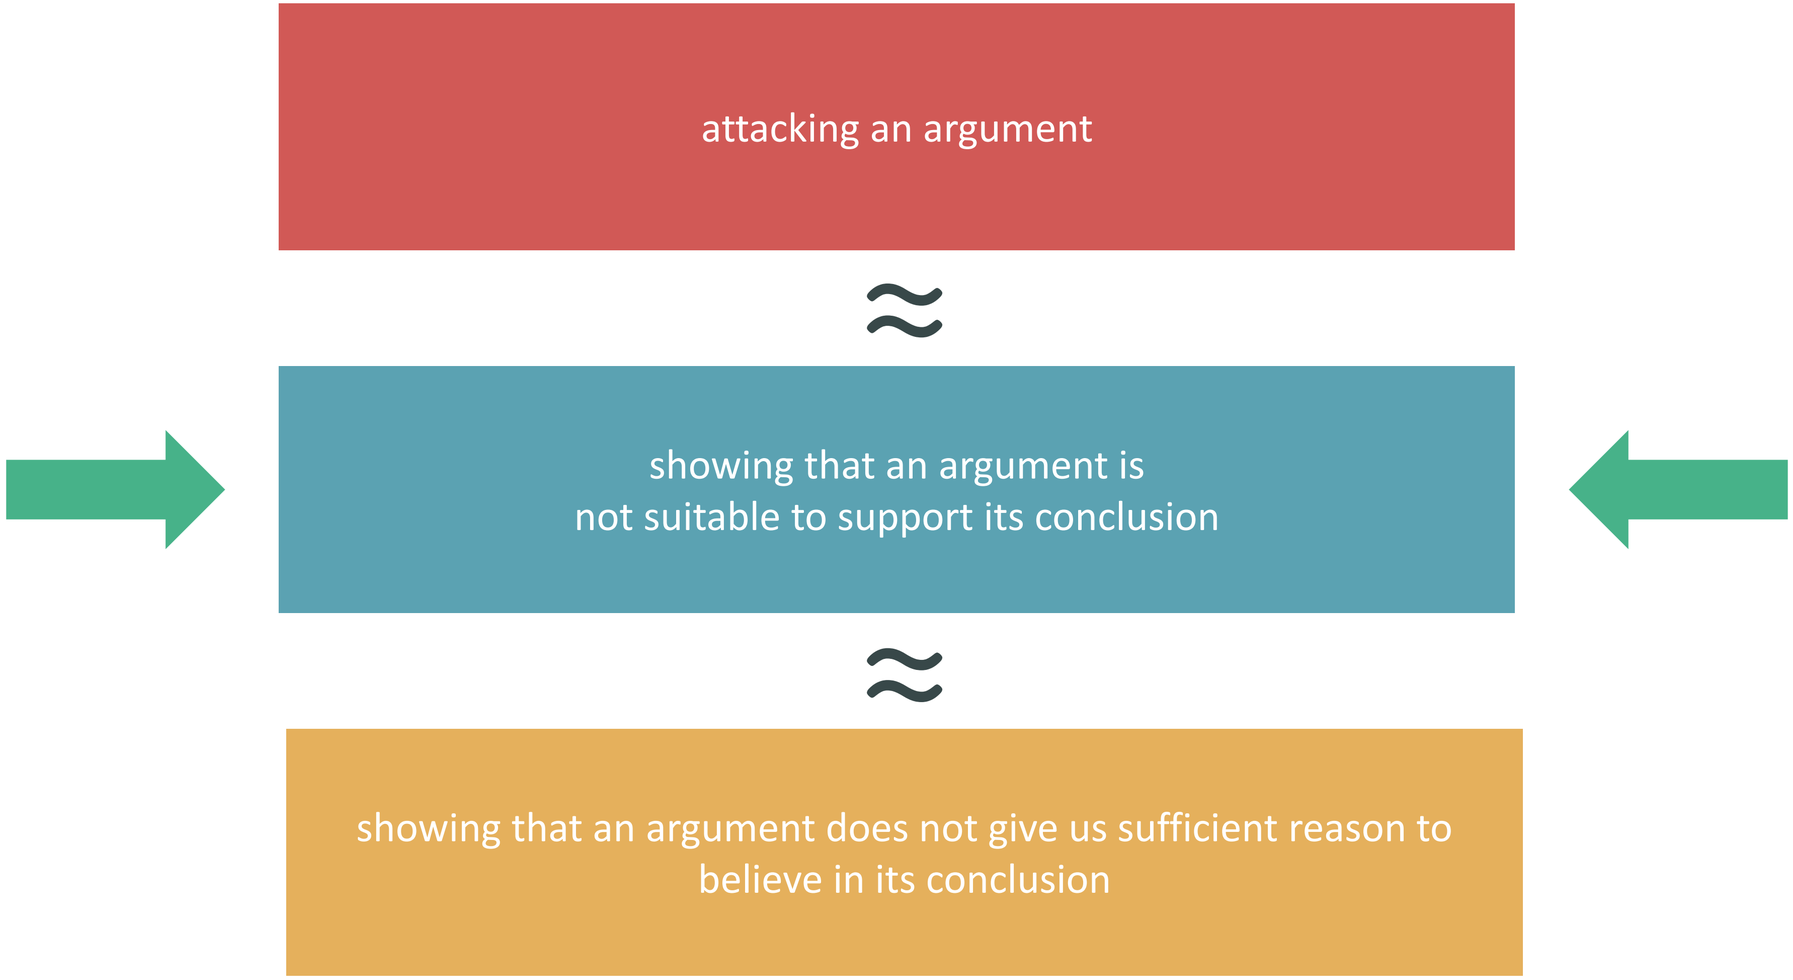

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

1

show that is has at least one false premise
(standard way)

2

show that the inference does not work

for deductive arguments:

show that the argument is not valid

for defeasible arguments:

show that the argument is not defeasibly forceful

3

show that the reasoning is fallacious or that it is pseudo-reasoning

4

for defeasible arguments:

show that there is a defeater

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

1. Show that is has at least one false premise (standard way)

P1

P2

P3

C

¬ P2

If ¬ P2, then the argument is not suitable to support C

the argument is not suitable to support C

Note: Usually, only an argument is given that has ¬P2 as a conclusion (and that is what you should do)

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

1. Show that is has at least one false premise (standard way)

Example:

P1: If the system is designed to be obscure, an attacker cannot know how to attack.

P2: If an attacker cannot know how to attack, then the system is secure.

C1: Therefore, if the system is designed obscure, then the system is secure.

P3: If C1, then Security through Obscurity is sufficient for securing a critical system.

C: Therefore, Security through Obscurity is sufficient for securing a critical system.

P1‘: An obscure design could be flawed, someone who knows the system could reveal details about it or an attacker could guess a vulnerability.

P2‘: If P1, then C.

C‘: It is not true that if a system is designed to be obscure, then an attacker cannot know how to attack.

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

2. Show that the inference does not work, i.e. that the argument is

- not valid (for deductive arguments)
- not defeasibly forceful (for defeasible arguments)

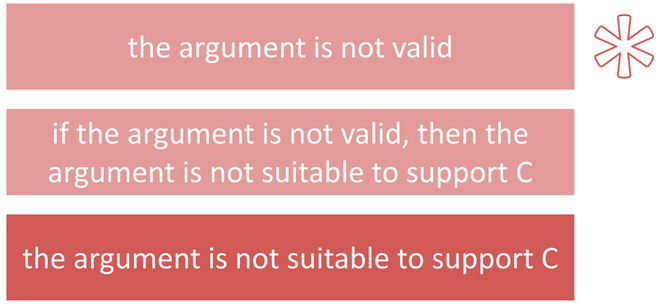

Note: Usually, only an argument in text form is given that shows why the argument is structurally flawed. (And that is what you should do.)

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

2. Show that the inference does not work

Example:

P1: If the system is completely obscure to an attacker, then an attacker cannot know how to attack.

P2: If attackers cannot know how to attack, then the system is secure.

P3: If a system that is designed to be obscure is secure, then Security through Obscurity is sufficient for securing a critical system.

C: Therefore, Security through Obscurity is sufficient for securing a critical system.

P1‘: The argument’s premises can be true, while the conclusion is false.

P2‘: If P1, then C.

C‘: The argument is not valid.

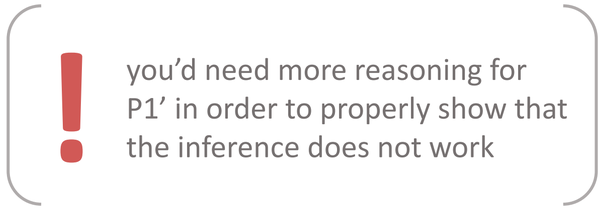

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

3. Show that the reasoning is fallacious or that it is pseudo-reasoning ⟶ **PT7 Fallacies and Pseudo Reasoning**

P1

P2

P3

C

The argument is a fallacy/pseudo-reasoning

if the argument is a fallacy/pseudo-reasoning, then the argument is not suitable to support C

the argument is not suitable to support C

Note: Usually, only an argument in text form is given that shows which kind of fallacy/pseudo reasoning the argument is and why. (And that is what you should do.)

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

3. Show that the reasoning is fallacious or that it is pseudo-reasoning ⟶ **PT7 Fallacies and Pseudo Reasoning**

Example:

P1: For critical systems, it suffices to use a design pattern called “Security though Obscurity” in order to make it secure.

P2: If P1, then no other measures but Security through Obscurity are needed to secure critical systems.

P3: If no other measures but Security through Obscurity are needed to secure critical systems, then Security by Obscurity is sufficient for securing a critical system.

C: Therefore, Security by Obscurity is sufficient for securing a critical system.

P1‘: The conclusion of the argument is presupposed in the first premise.

P2‘: If P1, then C.

C‘: The argument is a *petitio principii*.

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

4. Show that there is a defeater (for defeasible arguments)

P1

P2

P3

C

there is a defeater

if there is a defeater, then the argument is not suitable to support C

the argument is not suitable to support C

Note: Usually, only an argument is given that there is a defeater. (And that is what you should do.)

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion

4. Show that there is a defeater (for defeasible arguments)

Example:

P1: Obscurity makes it hard for an attacker to know where to attack.

P2: For most systems, it suffices for securing the system that it is hard for an attacker to know where to attack.

C: Probably, Security through Obscurity is sufficient for securing system X.

P1‘: System X is a critical system.

P2‘: If P1, then C.

C‘: It does not suffice for securing the system that it is hard for an attacker to know where to attack.

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion (yet)

1

show that is has at least one false premise
(standard way)

1lite

show that at least one premise is implausible and thereby shift the burden of proof

2

show that the inference does not work

for deductive arguments:

show that the argument is not valid

for defeasible arguments:

show that the argument is not defeasibly forceful

3

show that the reasoning is fallacious or that it is pseudo-reasoning

4

for defeasible arguments:

show that there is a defeater

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion (yet)

1lite. Show that at least one premise is implausible and thereby shift the burden of proof

P2 is implausible.

if P2 is implausible, then more reason for P2 needs to be given before this argument has to be taken as reason for C

more reason for P2 needs to be given before this argument has to be taken as reason for C

Note: Usually, you explicitly state that you try to shift the burden of proof, but the main argument will show why the premise in question is implausible. (And that is what you should do.)

! You do not show that the argument is unsuitable to support its conclusion, you only shift the burden of proof!

ATTACKING AN ARGUMENT

How to show that an argument is not suitable to support its conclusion (yet)

1lite. Show that at least one premise is implausible and thereby shift the burden of proof

Example:

P1: The cost of securing non-critical systems with other measures than Security through Obscurity is not outweighed by the benefits of doing so.

P2: If P1, then non-critical systems do not have to be secured with any other security measures than Security trough Obscurity.

C: Non-critical systems do not have to be secured with any other security measures than Security trough Obscurity.

P1‘: It is plausible that many security measures are not very costly and even the corruption of non-critical systems can be very dangerous if it happens on a large scale.

P2‘: If P1, then C.

C‘: It is implausible that the cost of securing non-critical systems with other methods than Security by Obscurity is not outweighed by the benefits of doing so.

… and this shifts the burden of proof.

ATTACKING AN ARGUMENT

How to show that an argument is not suitable *for a certain target audience* **to support its conclusion** *(yet)*

1

show that is has at least one false premise
(standard way)

1lite

show that at least one premise is implausible and thereby shift the burden of proof

2

show that the inference does not work

for deductive arguments:

show that the argument is not valid

for defeasible arguments:

show that the argument is not defeasibly forceful

3

show that the reasoning is fallacious or that it is pseudo-reasoning

4

for defeasible arguments:

show that there is a defeater

5

show that it is not rationally persuasive to the target audience

RATIONAL PERSUASIVENESS

Reminder: What is an argument for?

→ giving us good reason to believe in the conclusion

But when does an argument do this? When it is sound?

P1: If God exists, you should believe in God.
P2: God exists.

C: You should believe in God.

Suppose this argument is brought forward by a missionary to convince an atheist that she should believe in God. Should the atheist be convinced *by this argument*?

No, because from her point of view, she will have no reason to believe P2. Thus, the argument is not a good one (at least not if targeted at an atheist).

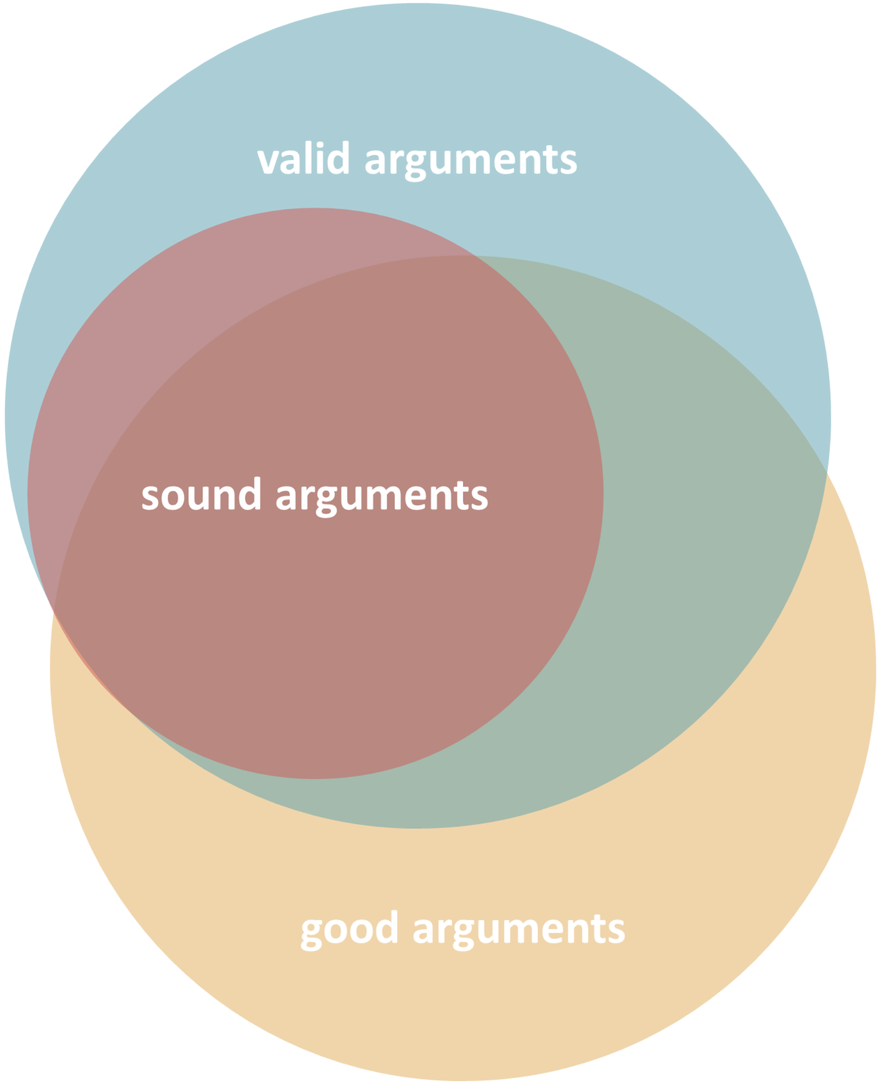

RATIONAL PERSUASIVENESS

General problem

sometimes, we do not know whether premises are true, i.e. we lack the means to decide whether the argument is sound.

(And this is actually true for most arguments in the lecture.)

But it also can be the case that it is very difficult or impossible

- for everyone (or for almost everyone) to decide whether a certain premise is plausible
- for a certain (target) audience whether a certain premise is plausible

RATIONAL PERSUASIVENESS

Example 1: Difficulties for everyone

P1: If the coin will land tails down, it will land heads up.

P2: The coin will land tails down.

C: The coin will land heads up.

It is generally difficult or even impossible to decide how a coin will land

Example 2: Difficulties for the target audience

P1: If the secret service does not do anything problematic without the knowledge of the general public, then you can trust the secret service.

P2: The secret service does not do anything problematic without the knowledge of the general public.

C: You can trust the secret service.

If you are not in the secret service, you cannot tell.

Such arguments are unpersuasive.

Rational Persuasiveness

An argument is rationally persuasive to someone iff it is (epistemically) rational for them to be persuaded by the conclusion given the argument.

When can this be the case?

- **valid argument** + they have sufficient evidence for the premises
- **defeasibly forceful argument** + they have sufficient evidence for the premises + they have sufficient evidence that the argument is not defeated

RATIONAL PERSUASIVENESS

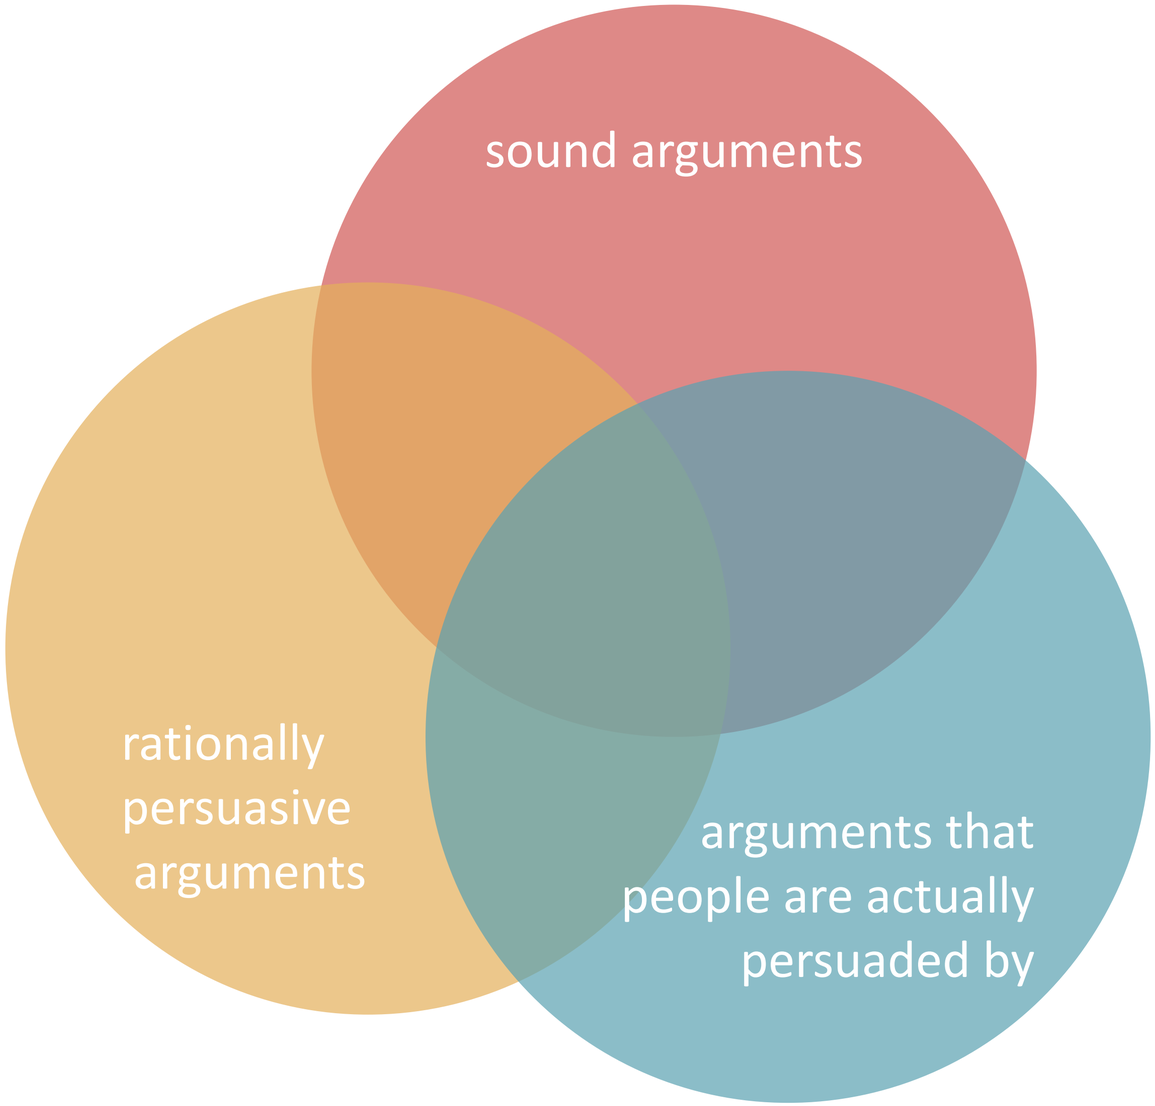

ATTACKING AN ARGUMENT

How to show that an argument is not suitable *for a certain target audience* **to support its conclusion** *(yet)*

1

show that is has at least one false premise
(standard way)

1lite

show that at least one premise is implausible and thereby shift the burden of proof

2

show that the inference does not work

for deductive arguments:

show that the argument is not valid

for defeasible arguments:

show that the argument is not defeasibly forceful

3

show that the reasoning is fallacious or that it is pseudo-reasoning

4

for defeasible arguments:

show that there is a defeater

5

show that it is not rationally persuasive to the target audience

ATTACKING AN ARGUMENT

How <u>NOT</u> to show that an argument is not suitable for any target audience ✓ **to support its conclusion** (ever)

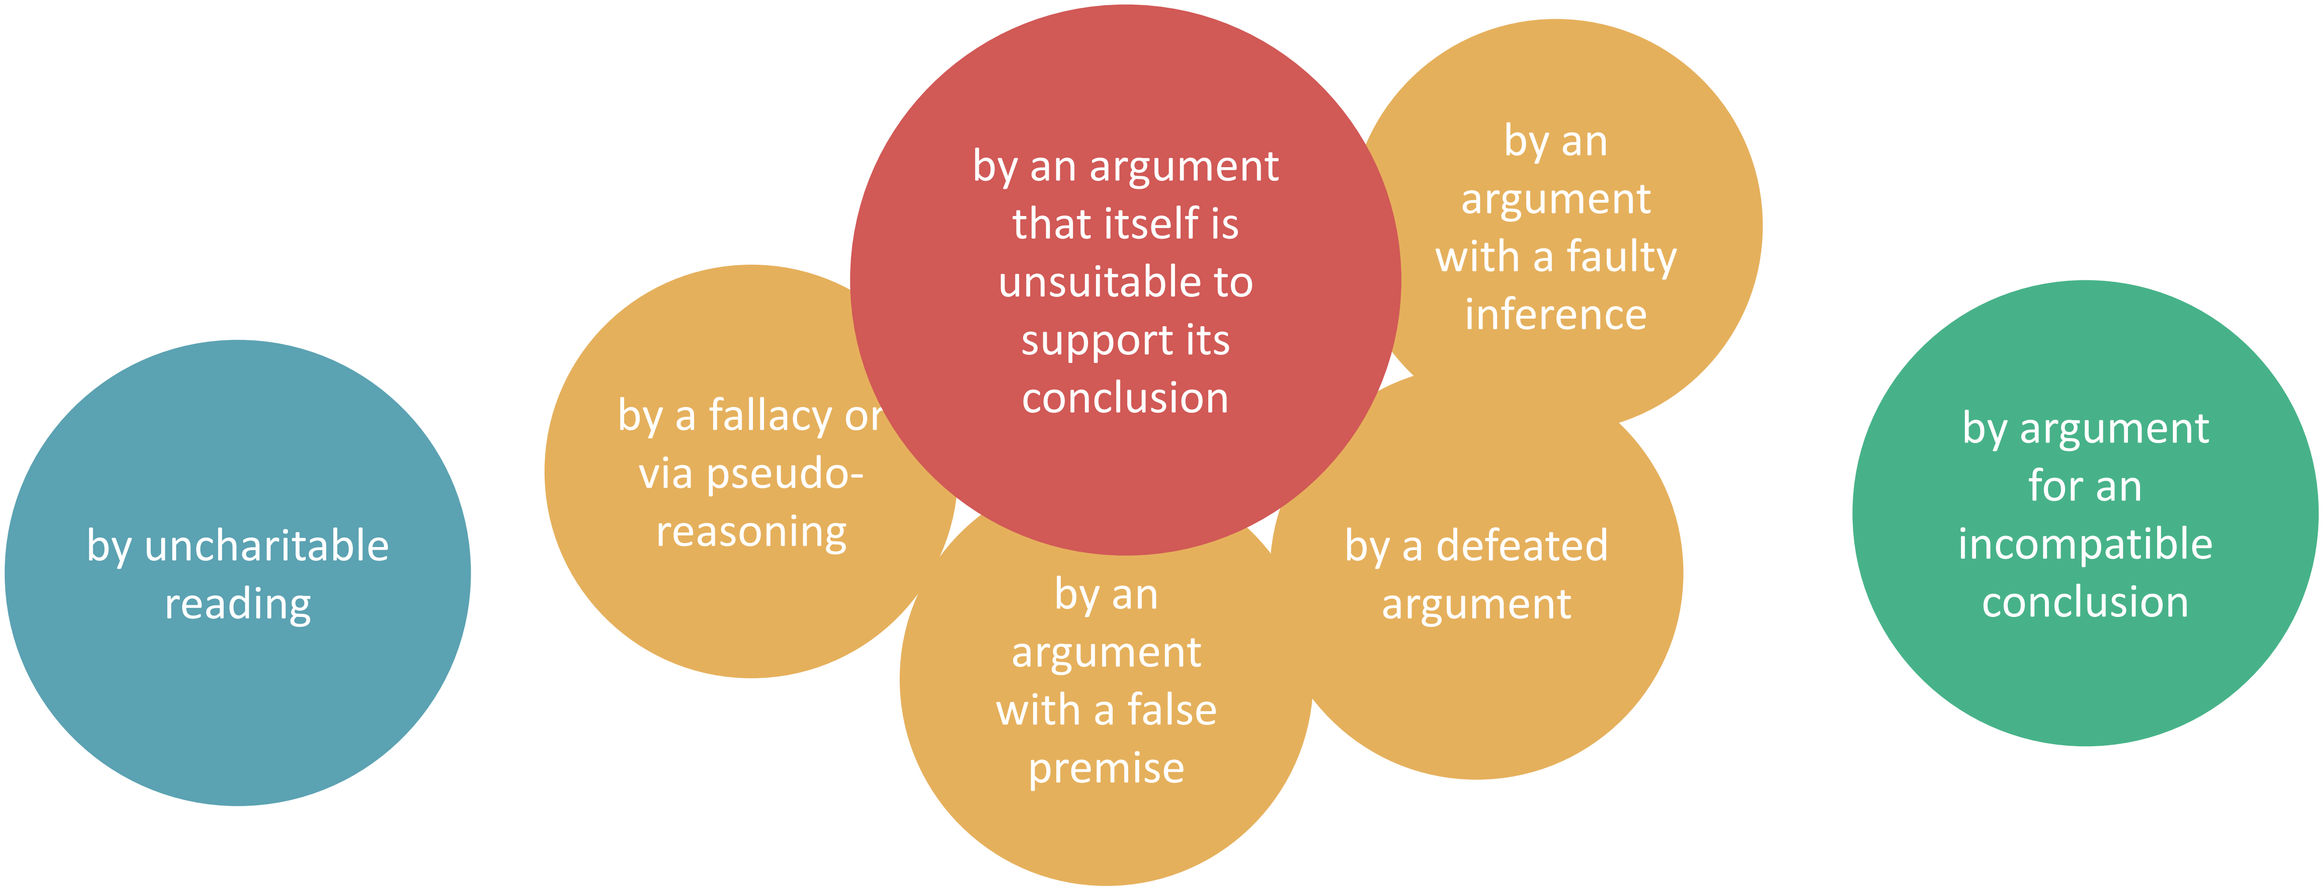

Ethics for Nerds

An Advanced Course in Computer Science
Summer Semester 2020

Precise Thinking 5.2
Discourse

The Structure of Proper Discourse

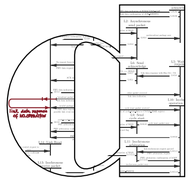

Prof. Holger Hermanns,
Kevin Baum, Sarah Sterz

UNIVERSITÄT
DES
SAARLANDES

PRECISE THINKING

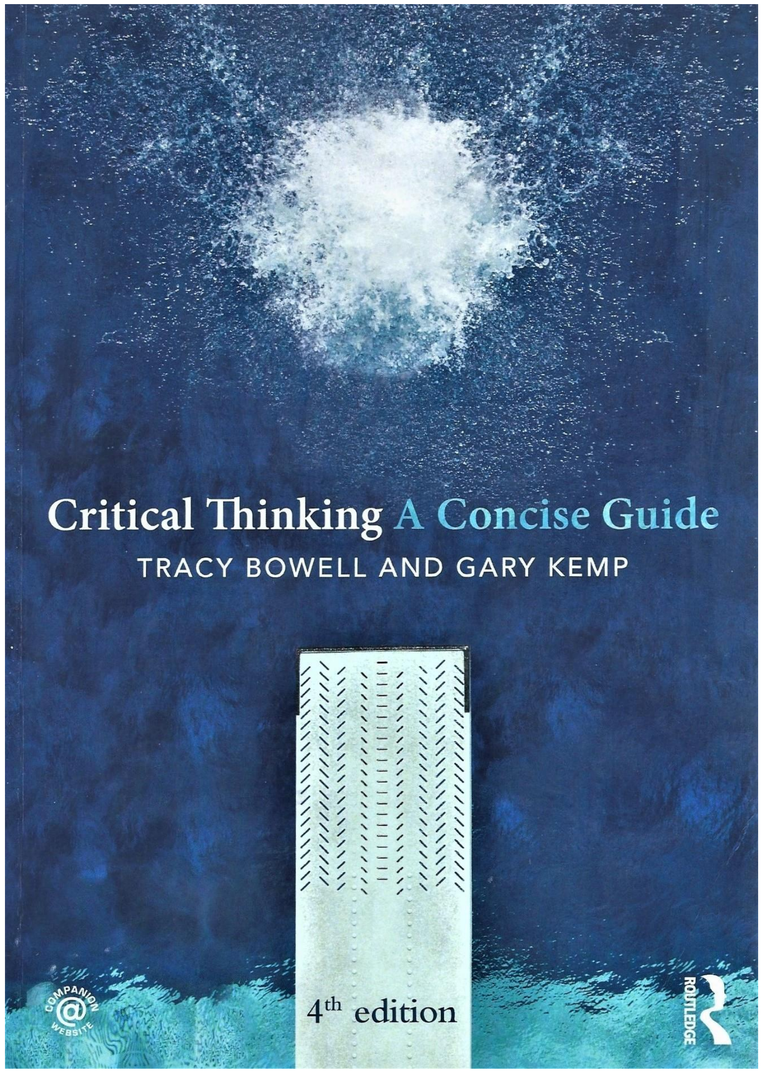

We *very* loosely follow this book:

> Bowell, T., & Kemp, G. (2015). *Critical thinking: A concise guide (4th ed.)*. Routledge.

However, *lots* of things deviate from the book. What is said in the lecture has precedence over the book.

You do **not** have to buy the book. If you want to have an inexpensive look, you can find a digital copy of the 2nd edition online at:

http://www.academia.edu/download/46383480/_Tracy_Bowell__Critical_Thinking_A_Concise_Guide_BookFi.org.pdf

DISCOURSE

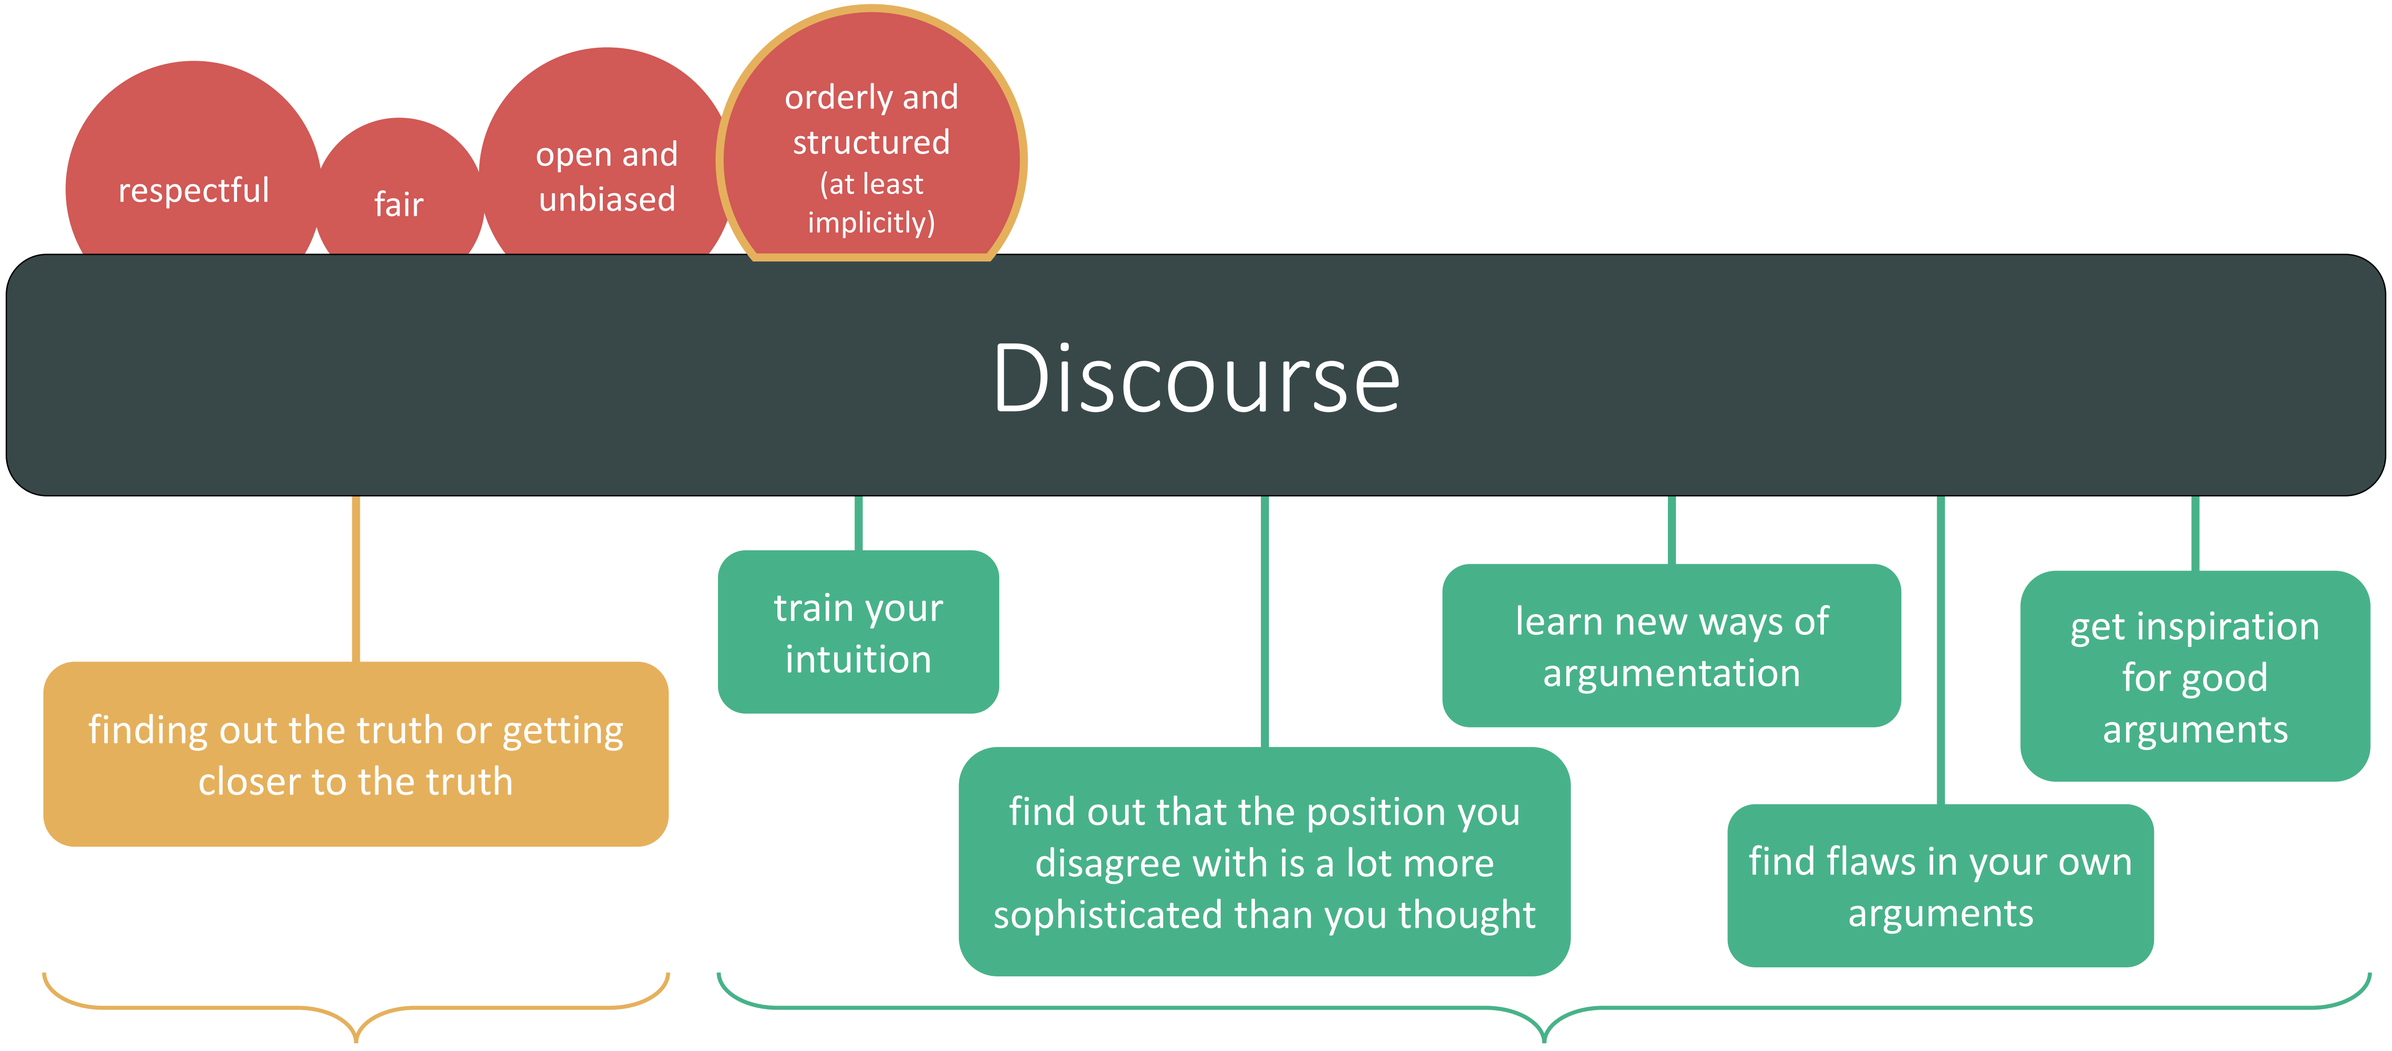
respectful
fair
open and unbiased
orderly and structured (at least implicitly)
Discourse
finding out the truth or getting closer to the truth
train your intuition
find out that the position you disagree with is a lot more sophisticated than you thought
learn new ways of argumentation
find flaws in your own arguments
get inspiration for good arguments
learn more about the world
improve your skills of how to think about the world

A BARE-BONES PROTOCOL FOR DISCUSSIONS

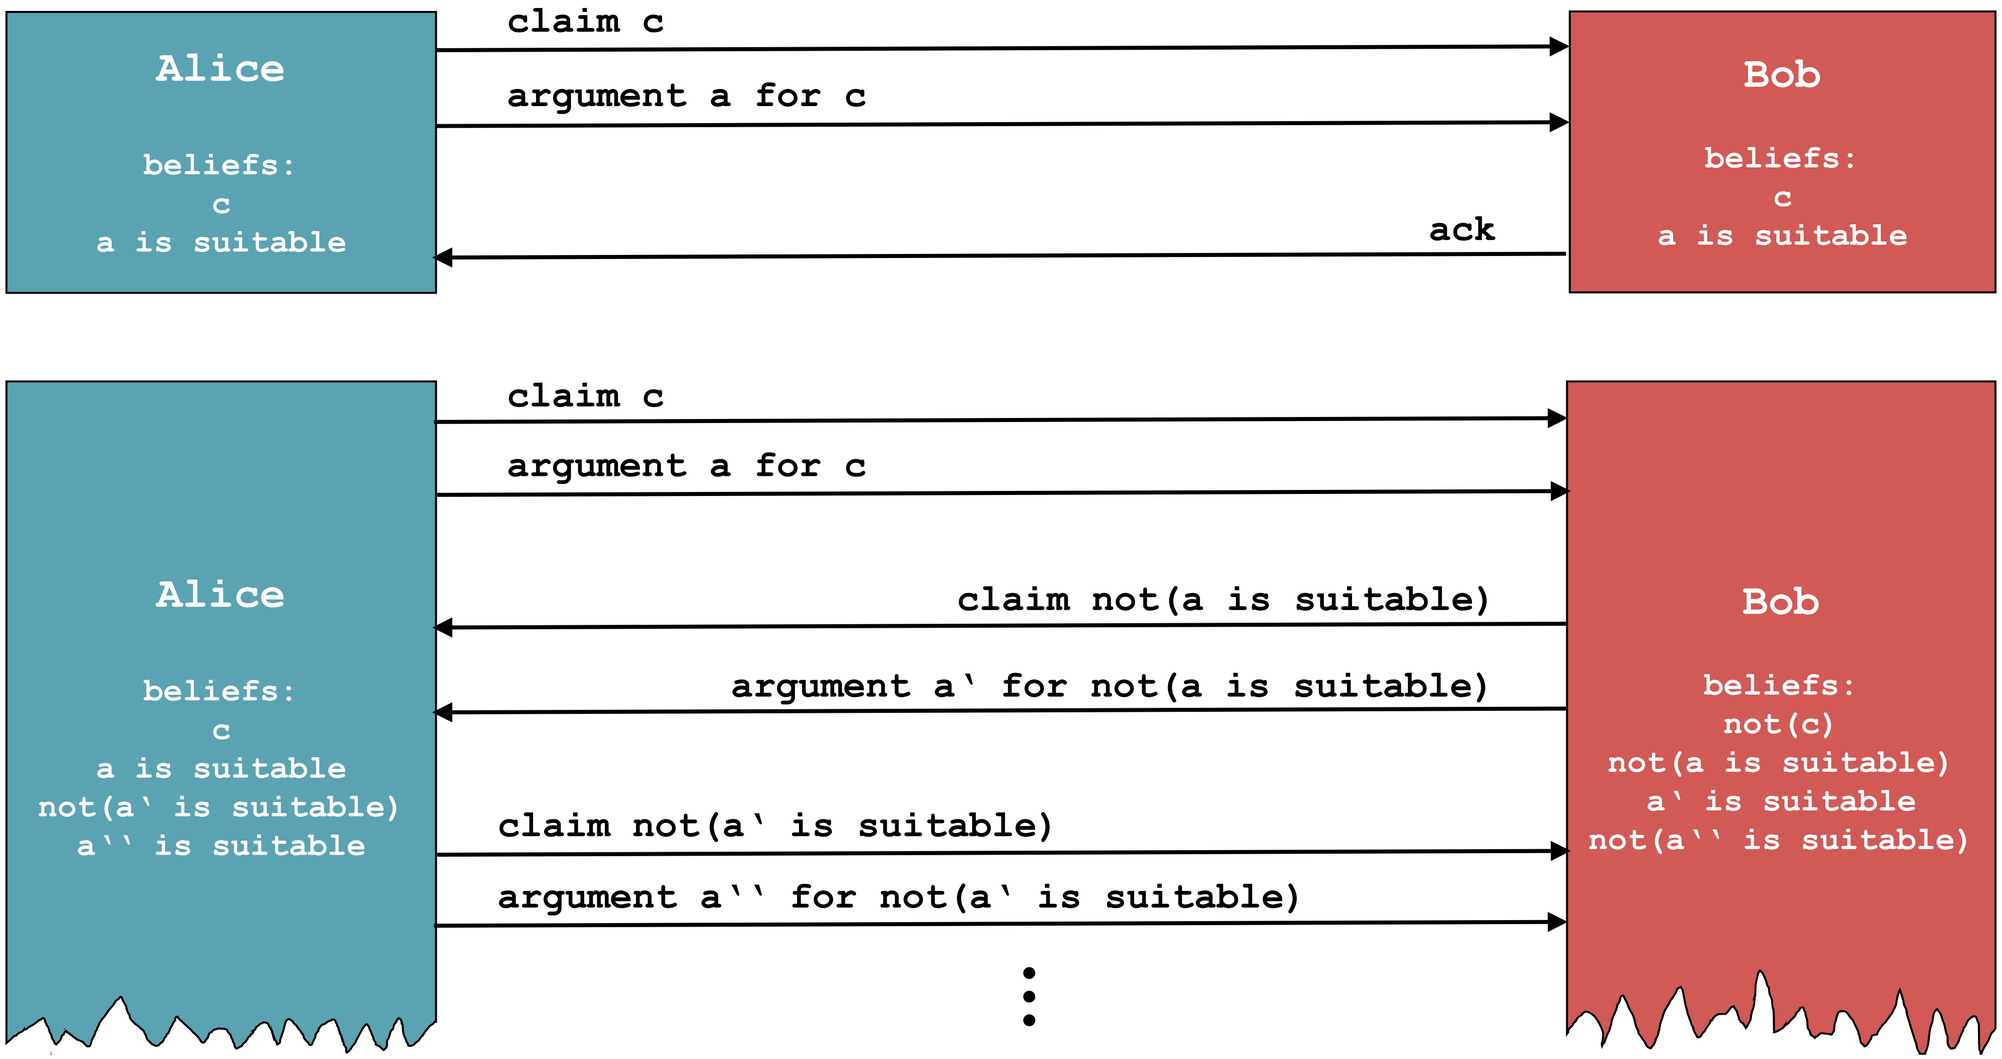

A BARE-BONES PROTOCOL FOR DISCUSSIONS

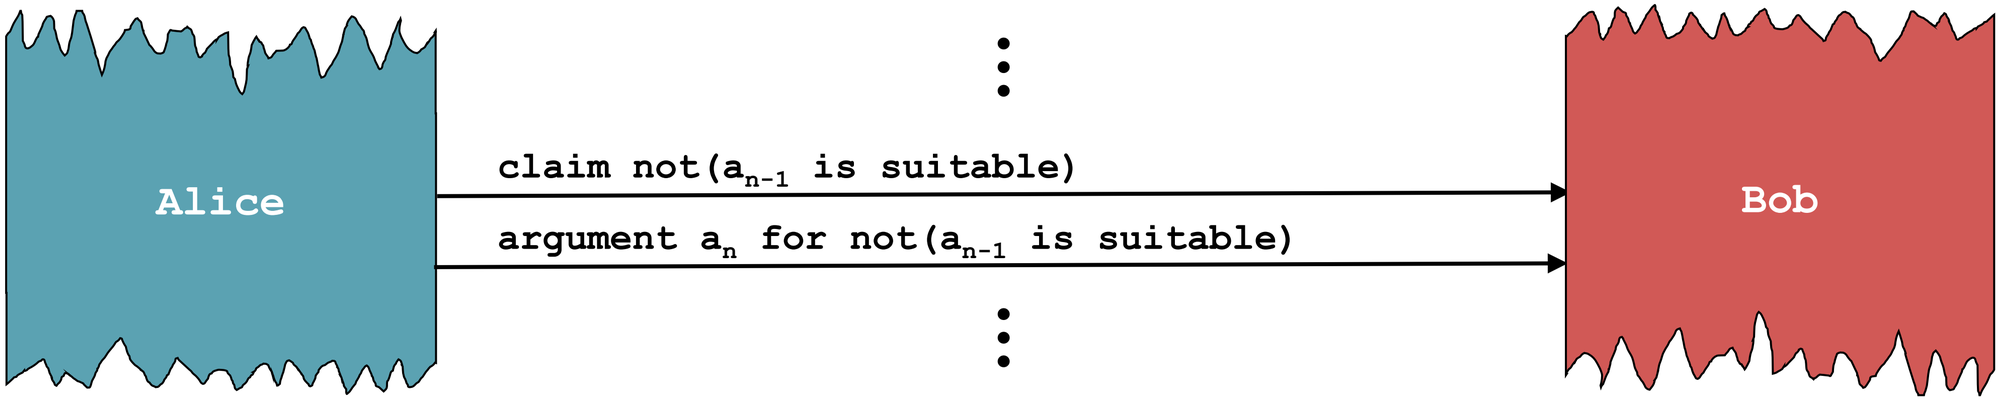

How does this process end?

- it continues forever
  (only theoretical possibility)
- they give up and agree to disagree
- they ultimately withhold judgement
  (or lower their credence in **c**)
- one of them is convinced by an argument $\mathbf{a_n}$ of the other person

Agreeing with another's argument

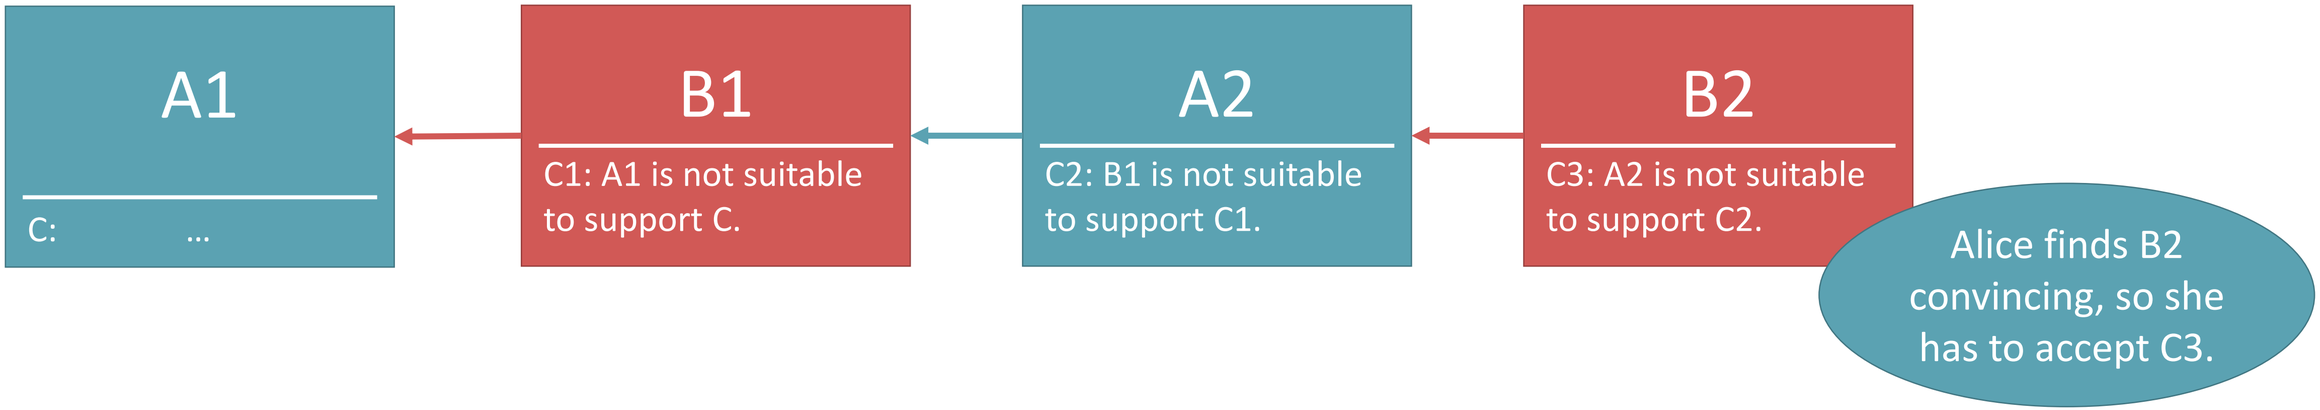

Agreeing with another's argument

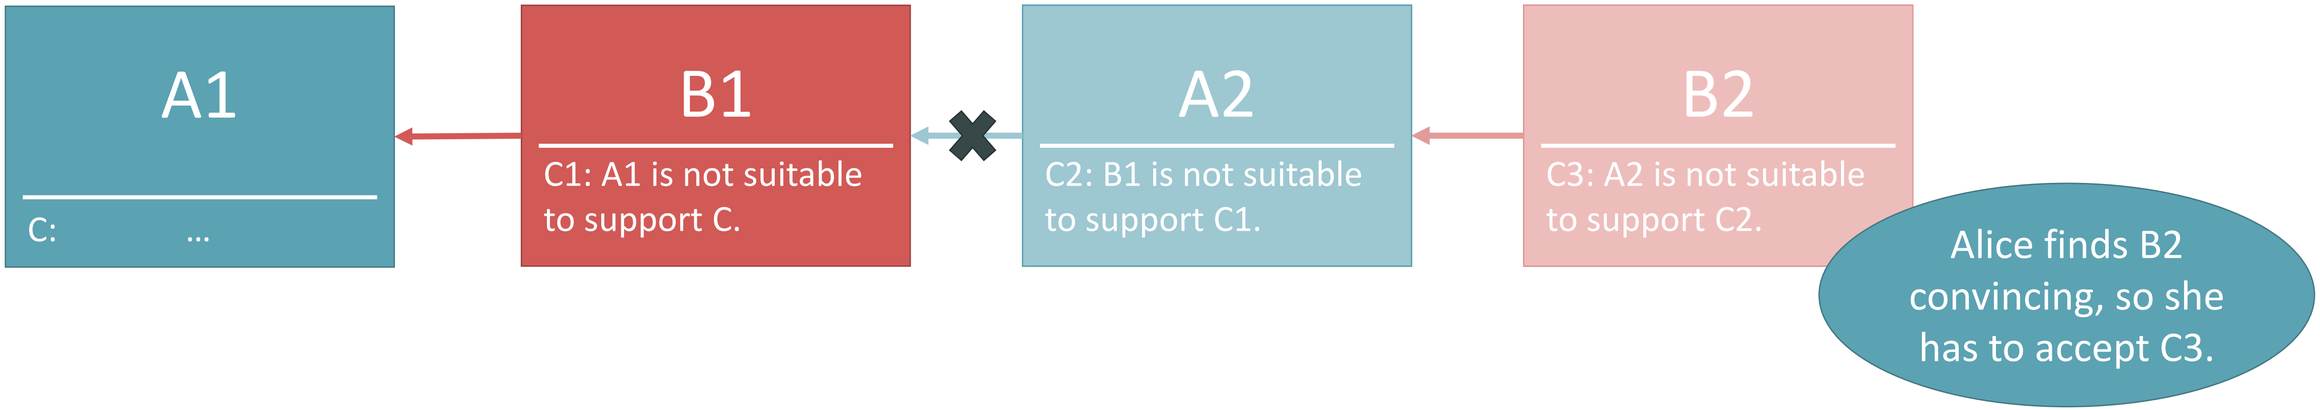

Agreeing with another's argument

Possibility 1: Alice comes up with a new argument against B1 that is robust against B2

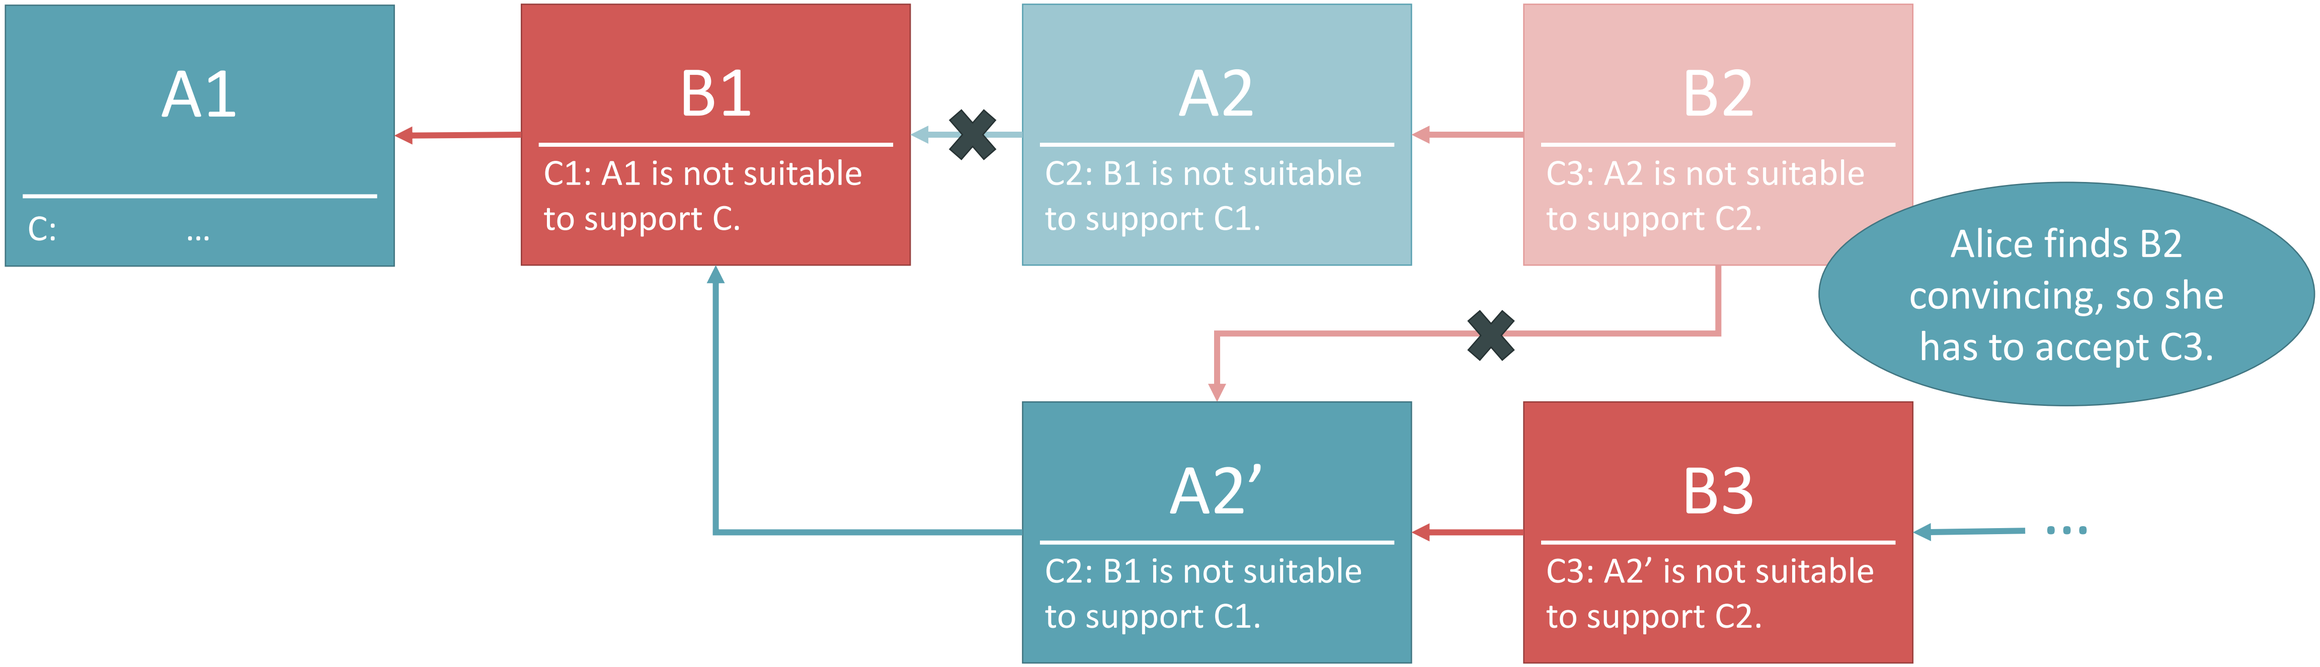

then the game resumes and Bob has to come up with a new argument that works against A2'

Agreeing with another's argument

Possibility 2: Alice cannot think of a new argument, so she needs to accept the conclusion of B1.

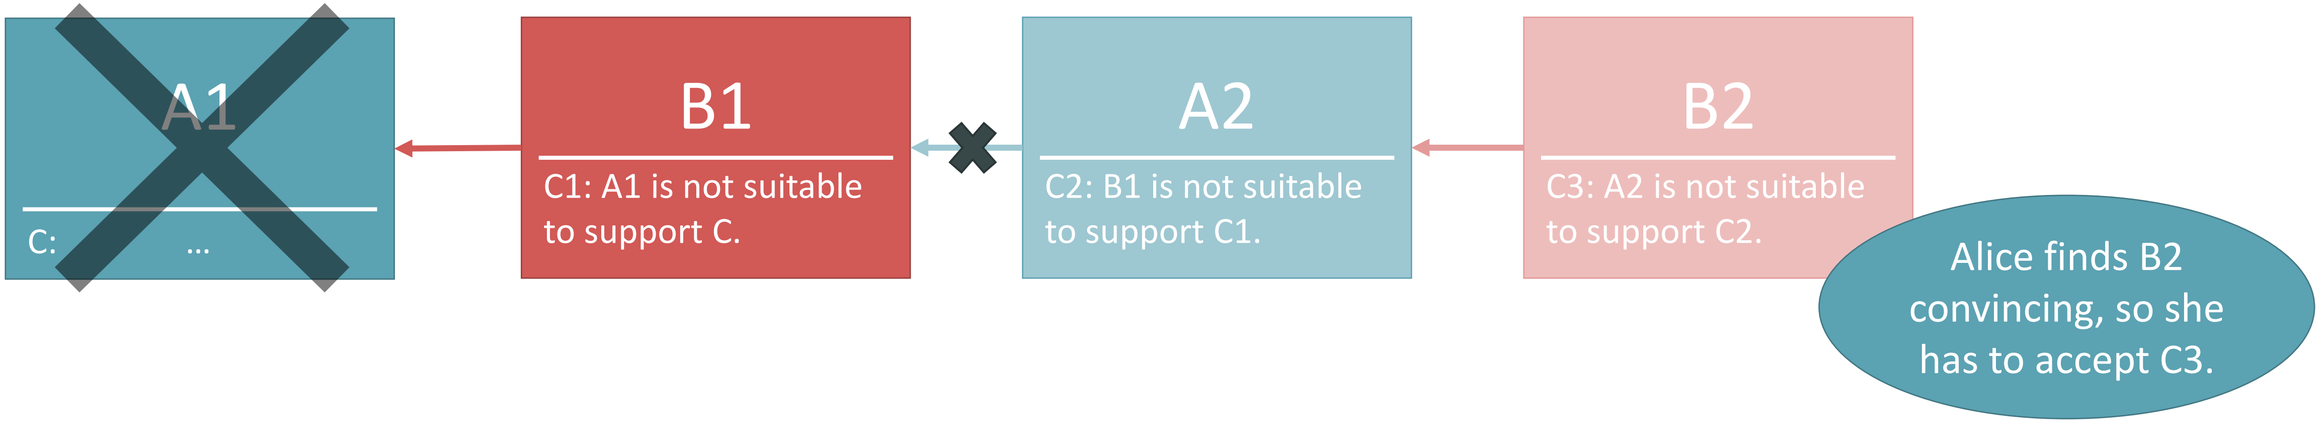

Then, she has to give up A1. Depending on how A1 and B1 work, she might only have to give up A1 and try to find a new argument for C, or even need to give up C entirely.

Example: Agreeing with another's argument

Argument A1

$P1_{A1}$: Ice cream is incredibly tasty.

$P2_{A1}$: If ice cream is incredibly tasty, it makes people happy.

$P3_{A1}$: If ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C: Therefore, unlimited access to ice cream should be a basic human right.

Example: Agreeing with another's argument

Argument A1

$P1_{A1}$: Ice cream is incredibly tasty.

$P2_{A1}$: If ice cream is incredibly tasty, it makes people happy.

$P3_{A1}$: If ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C: Therefore, unlimited access to ice cream should be a basic human right.

Argument B1

$P1_{B1}$: It is infeasible to give everyone unlimited access to ice cream.

$P2_{B1}$: People can live a life that is worthwhile without unlimited access to ice cream.

$P3_{B1}$: If $P1_{B1}$ and $P2_{B1}$, then it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C1: Therefore, it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

Example: Agreeing with another's argument

Argument A1

$P1_{A1}$: Ice cream is incredibly tasty.

$P2_{A1}$: If ice cream is incredibly tasty, it makes people happy.

$P3_{A1}$: If ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C: Therefore, unlimited access to ice cream should be a basic human right.

Argument B1

$P1_{B1}$: It is infeasible to give everyone unlimited access to ice cream.

$P2_{B1}$: People can live a life that is worthwhile without unlimited access to ice cream.

$P3_{B1}$: If $P1_{B1}$ and $P2_{B1}$, then it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C1: Therefore, it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

Argument A2

$P1_{A2}$: A human being needs unlimited access to ice cream for a life that is worth living.

$P2_{A2}$: If $P1_{A2}$, then people cannot live a life that is worthwhile without unlimited access to ice cream.

C2: Therefore, people cannot live a life that is worthwhile without unlimited access to ice cream.

Example: Agreeing with another's argument

Argument A1

$P1_{A1}$: Ice cream is incredibly tasty.

$P2_{A1}$: If ice cream is incredibly tasty, it makes people happy.

$P3_{A1}$: If ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C: Therefore, unlimited access to ice cream should be a basic human right.

Argument B1

$P1_{B1}$: It is infeasible to give everyone unlimited access to ice cream.

$P2_{B1}$: People can live a life that is worthwhile without unlimited access to ice cream.

$P3_{B1}$: If $P1_{B1}$ and $P2_{B1}$, then it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C1: Therefore, it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

Argument A2

$P1_{A2}$: A human being needs unlimited access to ice cream for a life that is worth living.

$P2_{A2}$: If $P1_{A2}$, then people cannot live a life that is worthwhile without unlimited access to ice cream.

C2: Therefore, people cannot live a life that is worthwhile without unlimited access to ice cream.

Argument B2

$P1_{B2}$: $P2_{A2}$ presupposes C2.

$P2_{B2}$: If $P1_{B2}$, then A2 is a petitio principii.

C3: Therefore, A2 is a petitio principii.

Example: Agreeing with another's argument

Argument A1

$P1_{A1}$: Ice cream is incredibly tasty.

$P2_{A1}$: If ice cream is incredibly tasty, it makes people happy.

$P3_{A1}$: If ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C: Therefore, unlimited access to ice cream should be a basic human right.

Argument B1

$P1_{B1}$: It is infeasible to give everyone unlimited access to ice cream.

$P2_{B1}$: People can live a life that is worthwhile without unlimited access to ice cream.

$P3_{B1}$: If $P1_{B1}$ and $P2_{B1}$, then it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C1: Therefore, it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

Argument A2

$P1_{A2}$: A human being needs unlimited access to ice cream for a life that is worth living.

$P2_{A2}$: If $P1_{A2}$, then people cannot live a life that is worthwhile without unlimited access to ice cream.

C2: Therefore, people cannot live a life that is worthwhile without unlimited access to ice cream.

Argument B2

$P1_{B2}$: $P2_{A2}$ presupposes C2.

$P2_{B2}$: If P1, then A2 is a petitio principii.

C3: Therefore, A2 is a petitio principii.

Alice finds B2 convincing, so she has to accept C3.

Example: Agreeing with another's argument – Possibility 1: Coming up with a new argument

Argument A1

$P1_{A1}$: Ice cream is incredibly tasty.

$P2_{A1}$: If ice cream is incredibly tasty, it makes people happy.

$P3_{A1}$: If ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C: Therefore, unlimited access to ice cream should be a basic human right.

Argument B1

$P1_{B1}$: It is infeasible to give everyone unlimited access to ice cream.

$P2_{B1}$: People can live a life that is worthwhile without unlimited access to ice cream.

$P3_{B1}$: If $P1_{B1}$ and $P2_{B1}$, then it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C1: Therefore, it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

Argument A2

$P1_{A2}$: A human being needs unlimited access to ice cream for a life that is worth living.

$P2_{A2}$: If $P1_{A2}$, then people cannot live a life that is worthwhile without unlimited access to ice cream.

C2: Therefore, people cannot live a life that is worthwhile without unlimited access to ice cream.

Argument B2

$P1_{B2}$: $P2_{A2}$ presupposes C2.

$P2_{B2}$: If P1, then A2 is a petitio principii.

C3: Therefore, A2 is a petitio principii.

Alice finds B2 convincing, so she has to accept C3.

Argument A2'

$P1_{A2'}$*: Unlimited access to ice cream is a basic human need.

$P2_{A2'}$: If $P1_{A2'}$, then people cannot live a life that is worthwhile without unlimited access to ice cream.

C2: Therefore, people cannot live a life that is worthwhile without unlimited access to ice cream.

Argument B3

$P1_{B3}$: Some humans don't even like ice cream.

$P2_{B3}$: Humans can survive without ice cream.

$P3_{B3}$: Ice cream is not an integral part of human life.

$P4_{B3}$: Societies existed for long times without ice cream.

$P2_{B3}$: If $P1_{B3}$ and $P2_{B3}$ and $P3_{B3}$ and $P4_{B3}$, then unlimited access to ice cream is not a basic human need.

C3: Therefore, unlimited access to ice cream is not a basic human need.

Example: Agreeing with another's argument – Possibility 2: Accept that A1 does not work

Argument A1

$P1_{A1}$: Ice cream is incredibly tasty.

$P2_{A1}$: If ice cream is incredibly tasty, it makes people happy.

$P3_{A1}$: If ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C: Therefore, unlimited access to ice cream should be a basic human right.

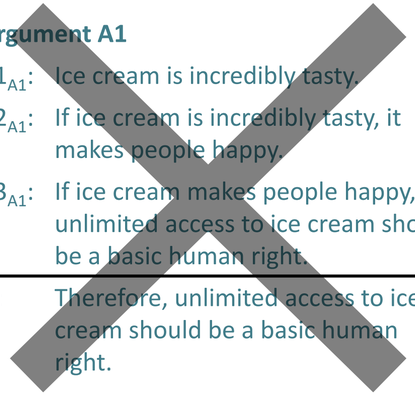

Argument B1

$P1_{B1}$: It is infeasible to give everyone unlimited access to ice cream.

$P2_{B1}$: People can live a life that is worthwhile without unlimited access to ice cream.

$P3_{B1}$: If $P1_{B1}$ and $P2_{B1}$, then it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

C1: Therefore, it is not true that if ice cream makes people happy, unlimited access to ice cream should be a basic human right.

Argument A2

$P1_{A2}$: A human being needs unlimited access to ice cream for a life that is worth living.

$P2_{A2}$: If $P1_{A2}$, then people cannot live a life that is worthwhile without unlimited access to ice cream.

C2: Therefore, people cannot live a life that is worthwhile without unlimited access to ice cream.

Argument B2

$P1_{B2}$: $P2_{A2}$ presupposes C2.

$P2_{B2}$: If P1, then A2 is a petitio principii.

C3: Therefore, A2 is a petitio principii.

Alice finds B2 convincing, so she has to accept C3.

A BARE-BONES PROTOCOL FOR DISCUSSIONS

Recap: How to **not** show that an argument is not suitable to support its conclusion:

by argument for an incompatible conclusion

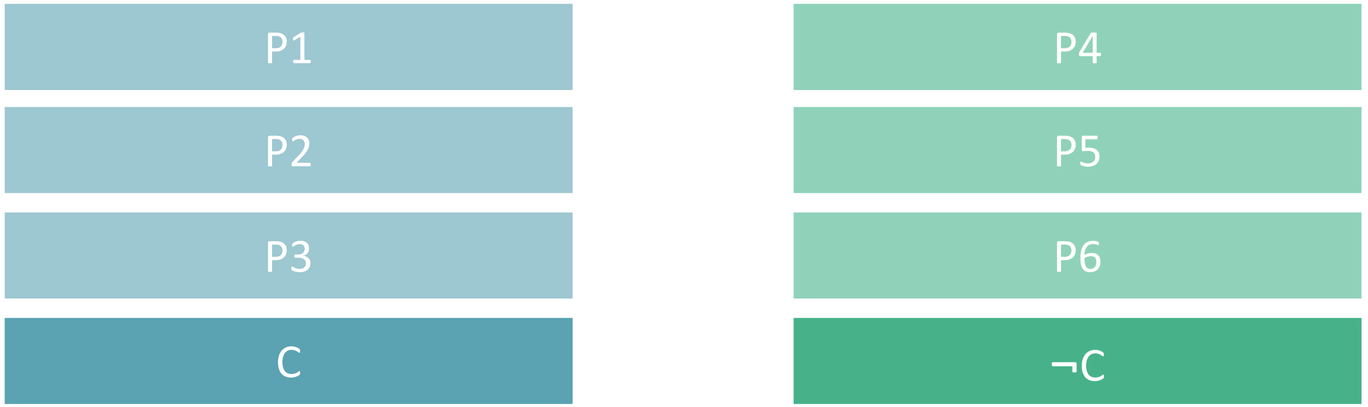

You always need to criticize an argument to refute it, not just give an argument for another conclusion

Note

In reality it often works like this:

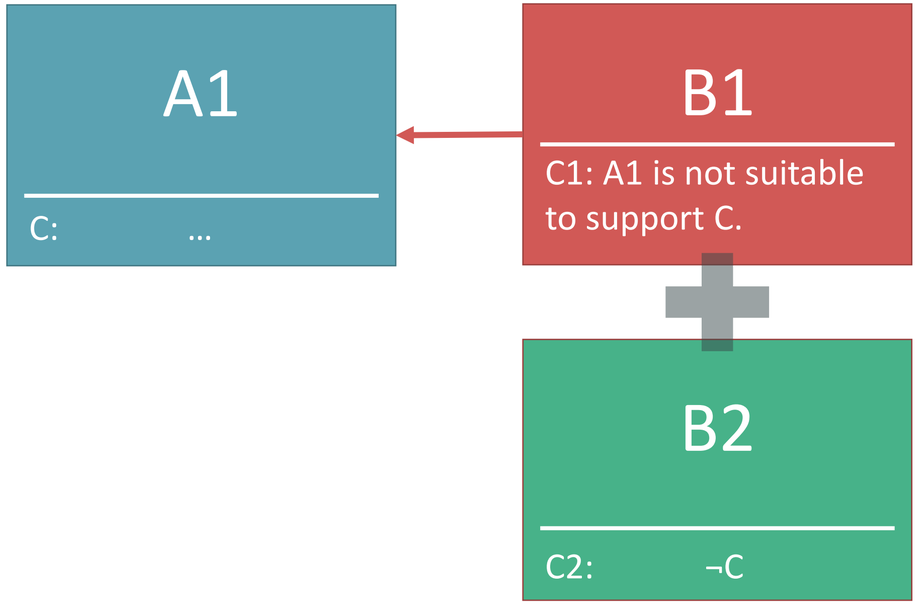

People try to refute the argument while at the same time giving an independent argument for the negation of the conclusion.

Side note: moving targets

The subject matter should not be changed and if it is changed, this has to be clearly indicated.

Otherwise, you are a moving target and this is a sin.

If argument is changed, everything from the discussion that builds upon this argument can become obsolete:

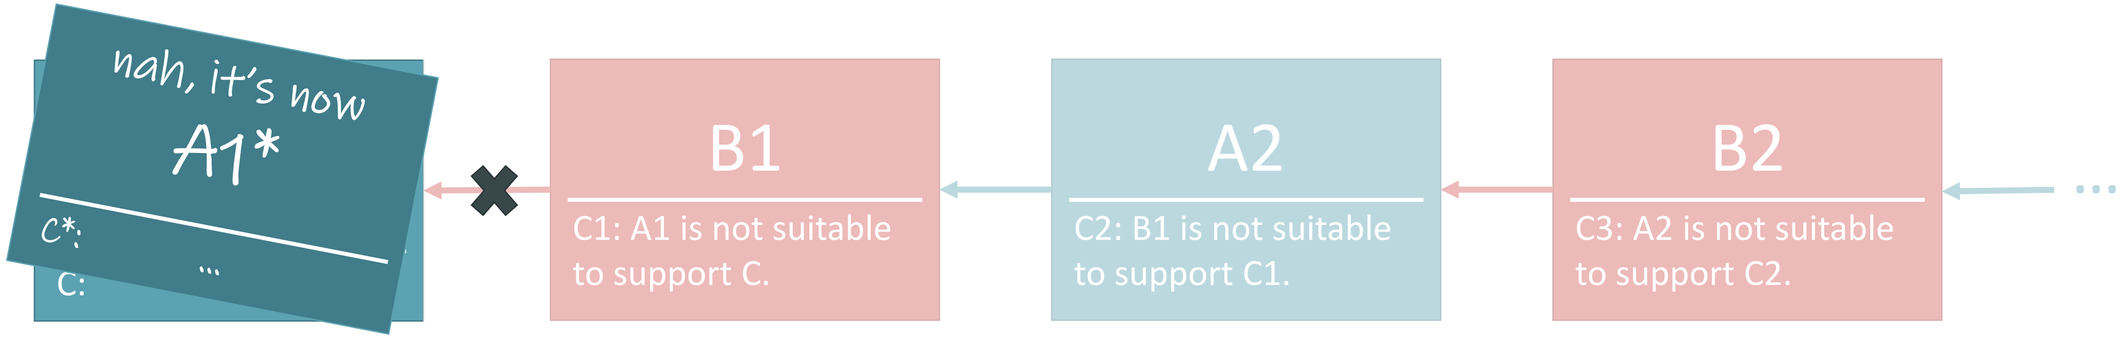

Which is why you should keep unnecessary changes to a minimum and at least notify the other of any changes to your arguments or your claims.

Ethics for Nerds

An Advanced Course in Computer Science
Summer Semester 2020

Precise Thinking 5.3
Discourse

Arguing about Morals

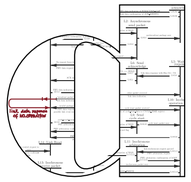

Prof. Holger Hermanns,
Kevin Baum, Sarah Sterz

UNIVERSITÄT
DES
SAARLANDES

Before you proceed:

Please watch the videos on ethics first

if you have not already done so!

ARGUMENTS ABOUT MORALS

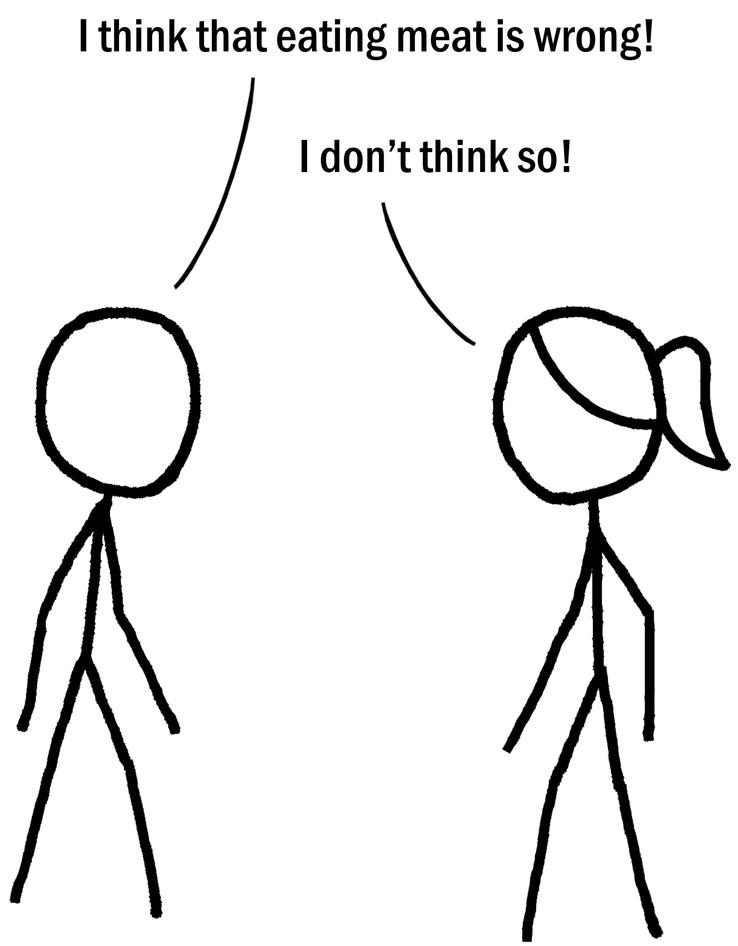

People can have disagreements about morals,

and disagreements call for argumentative discourse.

Four common kinds of arguments about morals:

reasoning from single moral theory

Reasoning from single moral theory

If you are committed to moral theory T and T says that φ-ing has a certain deontic/normative status, then you also are (or should be) committed to saying that φ-ing has this deontic/normative status.

P1: Moral theory T says that φ-ing is allowed for agent A in context C.

P2: Moral theory T is adequate for agent A in context C.

P3: If P1 and P2, then φ-ing is allowed for agent A in context C.

C: Therefore, φ-ing is allowed for agent A in context C.

What is behind this "adequate"?

	particular commitment to T	general commitment to T
no commitment to unique T	T is an adequate theory for agent A in context C	T is an adequate theory for all agents in all contexts
commitment to unique T	T is the adequate theory for agent A in context C	T is the adequate theory for all agents in all contexts

Example:

P1: Utilitarianism says that you ought to pull the lever in a trolley scenario.

P2: Utilitarianism is adequate in a trolley scenario.

P3: If P1 and P2, then you ought to pull the lever in a trolley scenario.

C: Therefore, you ought to pull the lever in a trolley scenario.

What is behind this "adequate"?

	particular commitment to Util.	general commitment to Util.
no commitment to unique Util.	Util. is an adequate theory for agents in trolley cases	Util. is an adequate theory for all agents in all contexts
commitment to unique Util.	Util. is the adequate theory for agents in trolley cases	Util. is the adequate theory for all agents in all contexts

Likely points of disagreement:

adequacy of moral theory (P2)

the right application of the moral theory (P1)

P1: Moral theory T says that φ-ing is allowed for agent A in context C.

P2: Moral theory T is adequate for agent A in context C.

P3: If P1 and P2, then φ-ing is allowed for agent A in context C.

C: Therefore, φ-ing is allowed for agent A in context C.

ARGUMENTS ABOUT MORALS

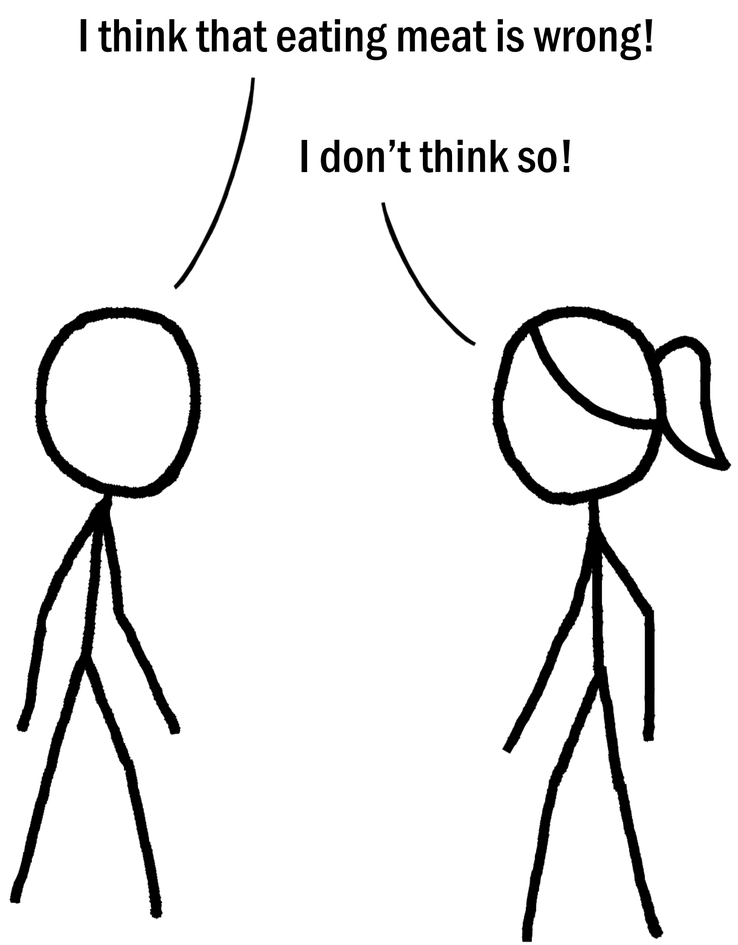

People can have disagreements about morals,

and disagreements call for argumentative discourse.

Four common kinds of arguments about morals:

reasoning from single moral theory

reasoning from a broad range of moral theories

Reasoning from broad range of moral theories

The shooting-a-shotgun-at-a-target technique: you do not commit to a moral theory, but when lots of very diverse, plausible moral theories from a set $\mathcal{T}$ tell you the same, it might be true.

P1:	All moral theories T $\in \mathcal{T}$ says that φ-ing is allowed for agent A in context C.
C:	Probably, φ-ing is allowed for agent A in context C.

The defeasible force of this argument increases with

- a greater number of moral theories
- more diverse moral theories
- more general moral theories
  (e.g. most consequentialist theories would be better than just equalizing, hedonistic, objective utilitarianism)

ARGUMENTS ABOUT MORALS

reasoning from a broad range of moral theories

Example:

P1:	Utilitarianism says that you should not kill without a good reason.
P2:	Kantianism says that you should not kill without a good reason.
P3:	Scanlonianism says that you should not kill without a good reason.
C:	Probably, you should not kill without good reason.

reasoning from a broad range of moral theories

Likely points of disagreement:

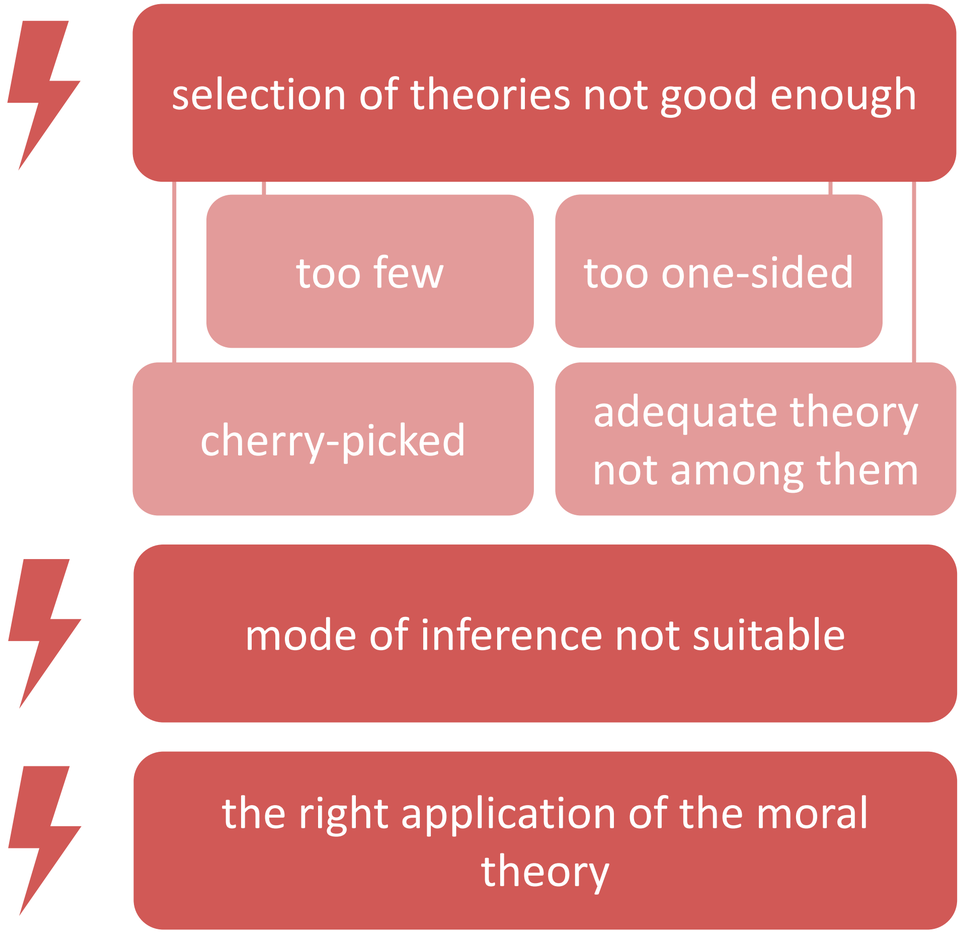

P1: All moral theories T $\in \mathcal{T}$ says that φ-ing is allowed for agent A in context C.

C: Probably, φ-ing is allowed for agent A in context C.

ARGUMENTS ABOUT MORALS

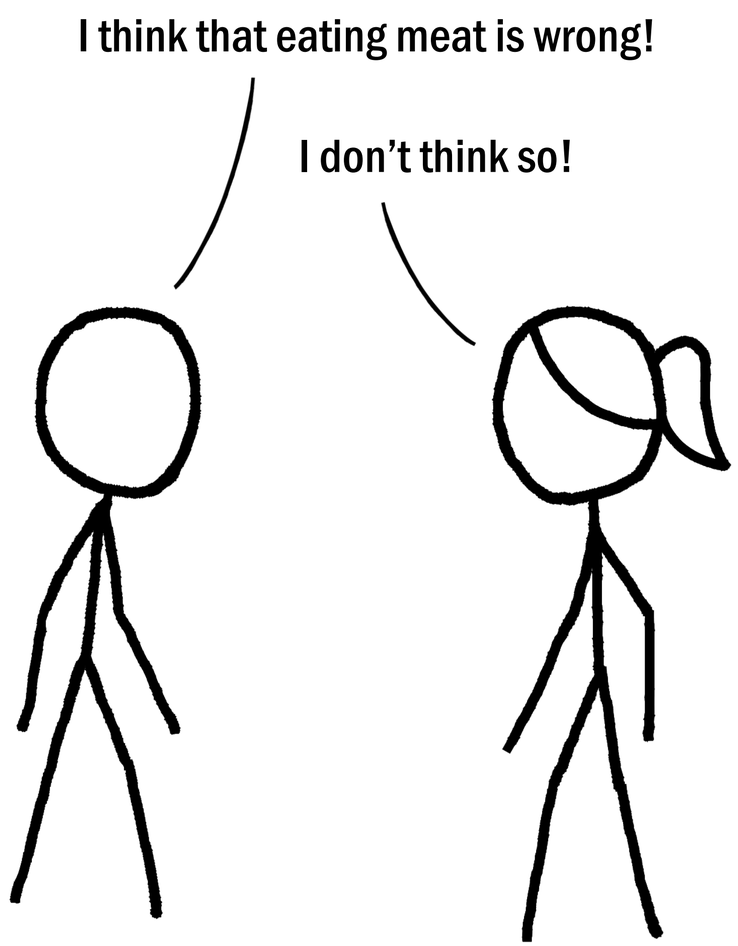

People can have disagreements about morals,

and disagreements call for argumentative discourse.

Four common kinds of arguments about morals:

- reasoning from single moral theory
- reasoning from a broad range of moral theories
- reasoning from consistent intuition

Reasoning from consistent intuition

You have a strong moral intuition that is not directly about φ-ing, but from that you can reason to the deontic status of φ-ing, because not every belief about the deontic status of φ-ing would be consistent with your strong intuition (given some other strong believes).

P1:	I/we have the strong intuition I and do not want to give it up.
P2:	If P1, then I/we should have intuition I*.
P3:	If I/we should have intuition I*, then I/we should accept that φ-ing is allowed for agent A in context C.
C:	Therefore, I/we should accept that φ-ing is allowed for agent A in context C.

Example:

P1: I/we have the strong intuition you ought to push the fat man in the fat man trolley scenario.

P2: If P1, then I/we should have the intuition that you ought to pull the lever in the regular trolley scenario.

P3: If I/we should have the intuition that you ought to pull the lever in the regular trolley scenario, then we should accept that you ought to pull the lever in the regular trolley scenario.

C: Therefore, I/we should accept that you ought to pull the lever in the regular trolley scenario.

Likely points of disagreement:

- we do not share intuition I, or are willing to give it up (P1)
- if one has intuition I, one does not need to have I* (P2)
- from intuition I*, nothing follows about φ-ing (P3)

P1: I/we have the strong intuition I and do not want to give it up.

P2: If P1, then I/we should have intuition I*.

P3: If I/we should have intuition I*, then we should accept that φ-ing is allowed for agent A in context C.

C: Therefore, I/we should accept that φ-ing is allowed for agent A in context C.

ARGUMENTS ABOUT MORALS

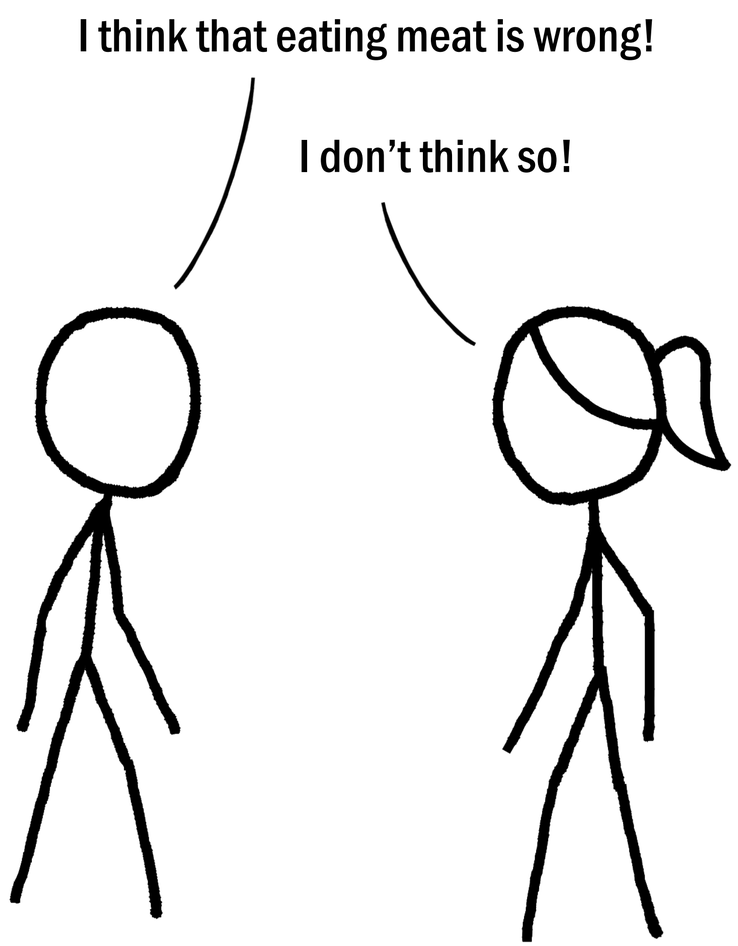

People can have disagreements about morals,

and disagreements call for argumentative discourse.

Four common kinds of arguments about morals:

- reasoning from single moral theory
- reasoning from a broad range of moral theories
- reasoning from consistent intuition
- reasoning from strong intuition

Reasoning from strong intuition

You just have a strong intuition about φ-ing.

P1:	I/we have the strong intuition that φ-ing is allowed for agent A in context C, and do not want to give it up.
P2:	If P1, then I/we should accept that φ-ing is allowed for agent A in context C.
C:	Therefore, I/we should accept that φ-ing is allowed for agent A in context C.

Example:

P1:	I/we have the strong intuition that inflicting needless pain on babies for fun is wrong, and I/we do not want to give that up.
P2:	If P1, then I/we should accept that inflicting needless pain on babies for fun is wrong.
C:	Therefore, I/we should accept that inflicting needless pain on babies for fun is wrong.

Likely point of disagreement:

we do not share intuition I, or are willing to give it up, or it is not strong (P1)

From intuition I*, nothing follows about φ-ing (P2)

P1: I/we have the strong intuition that φ-ing is allowed for agent A in context C, and do not want to give it up.

P2: If P1, then I/we should accept that φ-ing is allowed for agent A in context C.

C: Therefore, I/we should accept that φ-ing is allowed for agent A in context C.

ARGUMENTS ABOUT MORALS

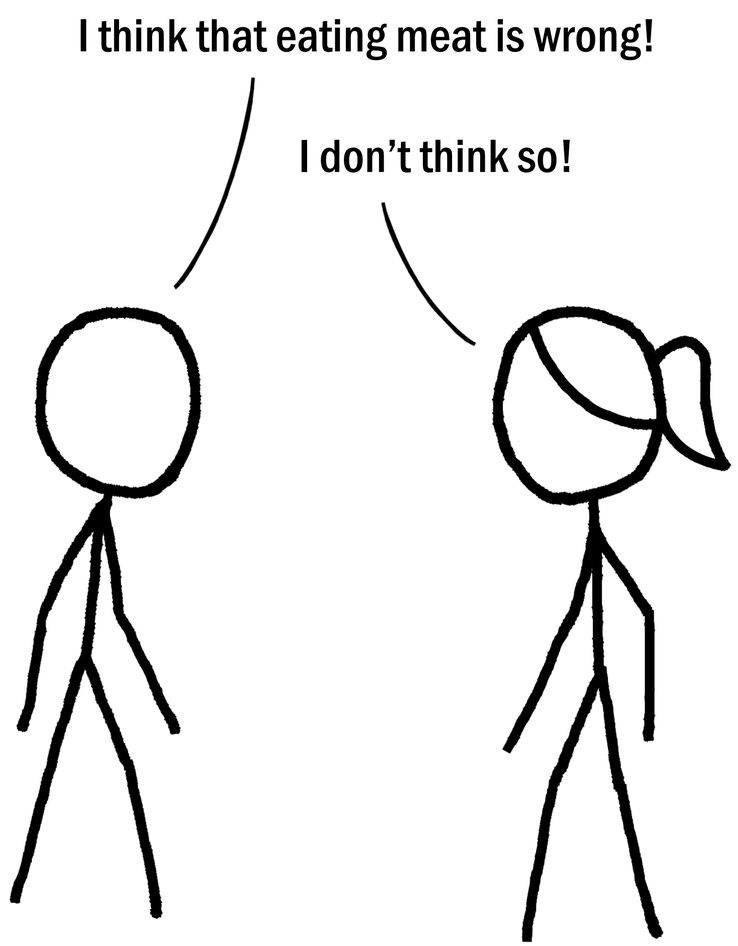

People can have disagreements about morals,

and disagreements call for argumentative discourse.

Four common kinds of arguments about morals:

- reasoning from single moral theory
- reasoning from a broad range of moral theories
- reasoning from consistent intuition
- reasoning from strong intuition